

# Contiuous Control Project

## Introduction

In this project, I solved the [Reacher](#) environment with a single agent. I used DDPG (Deep Deterministic Policy Gradient) algorithm.

## [Reacher environment](#)<sup>1,2,3</sup>

In this environment, a double-joint arm can move to target locations. The goal is that the agent must move its hand to the target location and needs to maintain its position at the target location as many time steps as possible. The agent reward function is +0.1 each step agent's hand is in target location. The observation space consists of 33 variables corresponding to position, rotation, velocity, and angular velocities of the arm. Each action is a vector with four numbers, corresponding to torque applicable to two joints. Every entry in the action vector should be a number between -1 and 1.

## [DDPG \(Deep Deterministic Policy Gradient\)](#)<sup>1,3</sup>

DDPG is an algorithm that simultaneously learns a Q-function and a policy. It uses off-policy data and the Bellman equation to learn the Q-function and uses the Q-function to learn the policy.

## Goal

The agent must get an average score of +30 over 100 consecutive episodes.

## Learning Algorithm

I followed the code provided by Udacity repository DDPG Pendulum.<sup>1</sup> There are some experimentation that I did to improve the performance of the agent.

1. I changed the nodes of the Actor and Critic fully connected (fc) layers, such as 512/256, 128/64, 64/64, with 3 fc layers, 300/300/300, etc. I also tried asymmetric fc layers between actor and critic.<sup>4</sup> This attempt failed. The reward didn't increase even until 200 episodes.
2. I added batch normalization to the first fully connected layer both Actor & Critic.
3. I tried the sigma value of 0.1, 0.2, 0.3, 0.5 and added mu value to 0.01, 0.1, 0.2
4. I increased the batch size from 128 to 1028
5. I increased the max\_t from 300 to 10000

Until step 5, I still could not make my agent to learn. I had poor results around 1-2 until 200 episodes

Below are sample of results from some of my experimentation:

Episode 50	Average Score: 0.24	
Episode 51	Average Score: 0.24	Time Duration: 3.30
Episode 52	Average Score: 0.25	Time Duration: 3.34
Episode 53	Average Score: 0.25	Time Duration: 3.35
Episode 54	Average Score: 0.25	Time Duration: 3.41
Episode 55	Average Score: 0.26	Time Duration: 3.57

---

Episode 100	Average Score: 1.62
Episode 200	Average Score: 1.88
Episode 235	Average Score: 1.60

6. Then, I decreased the LR\_ACTOR and LR\_CRITIC from  $1e-3$  to  $2e-4$ , used fc layers (both actor and critic) 200/200, added batch normalization to the first fc. I was able to improve the performance of the agent, however, the reward went down after it reached average score of 27 at around 400+ episodes.

Episode 100	Average Score: 2.56
Episode 200	Average Score: 4.91
Episode 300	Average Score: 15.11
Episode 400	Average Score: 25.32
Episode 500	Average Score: 25.05
Episode 600	Average Score: 23.87
Episode 660	Average Score: 20.69

7. With the help of a mentor, he mentioned that changing the Actor and Critic learning rates, together with the buffer size, will improve the agent performance.  
I increased the buffer size from  $1e5$  to  $2e5$

Episode 100	Average Score: 2.70
Episode 200	Average Score: 5.54
Episode 300	Average Score: 8.36
Episode 400	Average Score: 12.43
Episode 500	Average Score: 26.20
Episode 541	Average Score: 30.17
Environment solved in 541 episodes	Average Score 30.17

8. By increasing the buffer size from  $2e5$  to  $3e5$ , I got better result.

Episode 100	Average Score: 2.71
Episode 200	Average Score: 8.71
Episode 300	Average Score: 17.67
Episode 400	Average Score: 27.51
Episode 428	Average Score: 30.06
Environment solved in 428 episodes	Average Score 30.06

#### Hyperparameters:

- BUFFER\_SIZE=  $3e5$   
The replay buffer size
- BATCH\_SIZE= 128  
Number of inputs processed per batch when running Stochastic Gradient Descent
- GAMMA= 0.99  
Discount factor of the Q-Learning Algorithm
- TAU:  $1e-3$   
To perform soft updates of the target network parameters
- LR\_ACTOR:  $2e-4$   
Learning rate of the actor
- LR\_CRITIC:  $2e-4$   
Learning rate of the critic


## Plot of Rewards

By increasing the Actor and Critic learning rates to  $2e-4$ , buffer size to  $3e5$ , reducing sigma value to 0.1, I solved the environment in 428 episodes.

```
Episode 100    Average Score: 2.71
Episode 200    Average Score: 8.71
Episode 300    Average Score: 17.67
Episode 400    Average Score: 27.51
Episode 428    Average Score: 30.06
Environment solved in 428 episodes    Average Score 30.06
```

## 7. Reward Plot

```
fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(np.arange(1, len(scores)+1), scores)
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.show()
```



## **Ideas for Future Work**

This is an interesting project as I need to find the right hyperparameters to adjust in order to get the expected result.

For the future projects:

- Make a video output of this project
- Try higher buffer size to see if I can solve the environment faster
- Do the 20 agents training
- Try different algorithm, PPO

## **References**

1. Deep Reinforcement Learning Nano Degree Udacity Course
2. <https://github.com/Unity-Technologies/ml-agents/blob/master/docs/Learning-Environment-Examples.md#reacher>
3. Lillicrap, TP., Hunt, JJ., Pritzel, A., Continuous Control with Deep Reinforcement Learning. <https://arxiv.org/pdf/1509.02971.pdf>
4. Pinto, L., Andrychowicz, M., Welinder, P., et al. Asymmetric Actor Critic for Image-Based Robot Learning  
<https://arxiv.org/pdf/1710.06542.pdf>